

From Eliza to XiaoIce: Challenges and Opportunities with Social Chatbots

Heung-Yeung Shum, Xiaodong He, Di Li

Microsoft Corporation

Abstract: Conversational systems have come a long way since their inception in the 1960s. After decades of research and development, we've seen progress from Eliza and Parry in the 60's and 70's, to task-completion systems as in the DARPA Communicator program in the 2000s, to intelligent personal assistants such as Siri in the 2010s, to today's social chatbots like XiaoIce. Social chatbots' appeal lies not only in their ability to respond to users' diverse requests, but also in being able to establish an emotional connection with users. The latter is done by satisfying users' need for communication, affection, as well as social belonging. To further the advancement and adoption of social chatbots, their design must focus on user engagement and take both intellectual quotient (IQ) and emotional quotient (EQ) into account. Users should want to engage with a social chatbot; as such, we define the success metric for social chatbots as conversation-turns per session (CPS). Using XiaoIce as an illustrative example, we discuss key technologies in building social chatbots from *core chat* to *visual awareness* to *skills*. We also show how XiaoIce can dynamically recognize emotion and engage the user throughout long conversations with appropriate interpersonal responses. As we become the first generation of humans ever living with AI, we have a responsibility to design social chatbots to be both useful and empathetic, so they will become ubiquitous and help society as a whole.

1. Introduction

One of the fundamental challenges in artificial intelligence (AI) is endowing the machine with the ability to converse with humans using natural language (Turing, 1950). Early conversational systems, such as Eliza (Weizenbaum, 1966), Parry (Colby, 1975), and Alice (Wallace, 2009), were designed to mimic human behavior in a text-based conversation, hence to pass the Turing Test (Turing, 1950; Shieber, 1994) within a controlled scope. Despite impressive successes, these systems, which were precursors to today's social chatbots, were mostly based on hand-crafted rules. As a result, they worked well only in constrained environments.

Since the 1990s, a lot of research has been done on task-completion conversational systems (Price, 1990; Hemphill et al., 1990; Dahl et al., 1994; Walk et al., 2001, 2002). Examples include systems for reserving airline tickets as in the DARPA Airline Travel Information System (ATIS) project and for travel planning as in the DARPA Communicator program. ATIS and Communicator systems are designed to understand natural language requests and perform a variety of specific tasks for users, such as retrieving flight information and providing information to tourists. Task-completion conversational systems have been typically based on data-driven, machine-learned approaches. Their performance is excellent only

within domains that have well-defined schemas (Glass et al., 1995, Walk et al., 2001, Raux et al., 2005; Andreani et al., 2006; Wang et al., 2011; Tur and Mori, 2011).

In the past several years, a tremendous amount of investment has been made to developing intelligent personal assistants (IPAs) such as Apple's *Siri*¹, Microsoft's *Cortana*², *Google Assistant*³, *Facebook M*⁴, and Amazon's *Alexa*⁵. These IPAs are often deployed on mobile devices and are designed to answer a wide range of questions. In addition to passively responding to user requests, they also proactively anticipate user needs and provide in-time assistance such as reminding of an upcoming event or recommending a useful service without receiving explicit requests from the user (Sarikaya 2017). The daunting challenge for such IPAs is that they must work well in many open domain scenarios as people learn to depend on them to manage their work and lives efficiently.

More recently, social chatbots, such as Microsoft's *XiaoIce*⁶, have emerged as a new kind of

¹ <https://www.apple.com/ios/siri/>

² <https://www.microsoft.com/en-us/cortana/>

³ <https://assistant.google.com/>

⁴ <https://developers.facebook.com/blog/post/2016/04/12/>

⁵ <https://developer.amazon.com/alexa/>

⁶ <https://www.msXiaoIce.com/>

conversational system made possible by significant progress in AI and wireless communications. The primary goal of a social chatbot is not necessarily to solve all the questions the users might have, but rather, to be a virtual companion to users. By establishing an emotional connection with users, social chatbots can better understand them and therefore help them over a long period of time. To communicate effectively, social chatbots interact with users through multiple modalities, including text, speech, and vision. Social chatbots and IPAs have become popular recently due to progress in many relevant perceptual and cognitive AI technologies, e.g., natural language understanding (Bengio et al., 2001; Mikolov et al., 2013; Sutskever et al., 2014; Bahdanau et al., 2015; Mesnil et al., 2013, 2015), speech recognition and synthesis (Hinton et al., 2012; Deng et al., 2013; Xiong et al., 2016; Qian et al., 2014; van den Oord et al., 2016), computer vision (Krizhevsky et al., 2012; He et al., 2016), information retrieval (Huang et al., 2013; Elkahky et al., 2015), multimodal intelligence (Fang et al., 2015; Vinyals et al., 2015; Karpathy and Fei-fei, 2015; He and Deng, 2017), and empathic conversational systems (Fung et al., 2016; Alam et al., 2017).

In the next section, we briefly review conversational systems developed over the past several decades and discuss their features and limitations. We then introduce the design principles of social chatbots where both *intellectual quotient* (IQ) and *emotional quotient* (EQ) must be considered. We describe key components in designing social chatbots including *core chat*, *visual awareness*, and *skills*. Using Microsoft’s XiaoIce, we show that social chatbots can engage in long conversations with users by understanding their needs, establishing an emotional connection, and offering support and assistance. We also share our vision for the evolution of chatbots in the future. We discuss that, with fundamental breakthroughs to occur on key open problems, the tremendous business opportunities they will make possible, and their potential impact on society as AI becomes an integral part of our lives.

2. Brief Survey of Conversational Systems

In this section, we review representative conversational systems that have been developed over five decades. Such systems include early

chatbots, task-completion systems, intelligent personal assistants, and social chatbots.

2.1 Eliza and Early Chatbots

Inspired by the *Turing test* proposed in 1950 (Turing, 1950), researchers and engineers have developed multiple conversational systems for chitchat (Weizenbaum, 1966; Colby, 1975, Shieber, 1994; Wallace 2009). These early chatbots are computer programs that conducted a conversation using audio or text. Such programs were often designed to convincingly simulate how a human would behave as a conversational partner; if they succeed, they pass the Turing test.

Human: Men are all alike.
Eliza: In what way?
Human: They're always bugging us about something specific or other.
Eliza: Can you think of a specific example?
Human: Well, my boyfriend made me come here.
Eliza: Your boyfriend made you come here?
Human: He says I'm depressed much of the time.
Eliza: I am sorry to hear you are depressed.
Human: It's true. I'm unhappy.
Eliza: Do you think coming here will help you not to be unhappy?
Human: I need some help, that much seems certain.
Eliza: What would it mean to you if you got some help?
Human: Perhaps I could learn to get along with my mother.
Eliza: Tell me more about your family.
...
Chat:

Figure 1. A conversation between Eliza and a human (Weizenbaum, 1966).

Eliza, created by Joseph Weizenbaum at MIT in 1966, is perhaps the very first chatbot known publicly. It can communicate with humans based on hand-crafted scripts (Weizenbaum, 1966). These scripts simulate a Rogerian psychotherapist and can only accept textual input. It does not understand a conversation, but rather searches for appropriate responses through pattern matching, combined with some intelligent phrasing. *Eliza* has only a limited scope of knowledge and can only chat with people within a constrained domain. Nevertheless, many users believed they were talking to a real person when *Eliza* first came out. Figure 1 shows a conversation between *Eliza* and a human.

Parry is another chatbot developed by Kenneth Colby (1975). It is designed to behave like a paranoid person. It passed the Turing test for the first time in history⁷. However, *Parry* is still rule-based and has a similar structure as *Eliza* but with better controlling structure, language understanding capabilities, and especially a mental model that can simulate the bot's emotions. For example, *Parry* will respond with hostility if the anger level is high.

Alice, or Artificial Linguistic Internet Computer Entity, was developed by Richard Wallace (2009) to allow users to customize their chatbots. It uses an Artificial Intelligence Markup Language (AIML), a derivative of XML; AIML has tags that allow bots to recursively call a pattern matcher so that the language can be simplified. *Alice* won the Loebner Prize⁸ three times in 2000, 2001, and 2004, an award for the most human-like system (Shieber, 1994). However, because of the limitations of AIML, the capacities of these chatbots are constrained. For example, *Alice* failed to pass the ultimate Turing Test⁹, partially because chitchat systems built using AIML cannot maintain a dialogue for a long period of time.

2.2 Task-completion Conversational Systems

In contrast to chitchat systems, task-completion systems are designed for accomplishing specific tasks. These systems usually operate on constrained domains (Glass et al., 1995; Seneff et al., 1998; Rudnicky et al., 1999; Levin et al., 2000; Walk et al., 2001, 2002; Raux et al., 2005; Andreani et al., 2006; Wang et al., 2011; Tur and Mori, 2011). Figure 2 illustrates the architecture of a traditional task-completion spoken dialog system.

This architecture comprises an automatic speech recognizer (ASR), a spoken language understanding (SLU) module, a dialog manager (DM), a natural language generator (NLG), and a text-to-speech (TTS) synthesizer. The ASR takes raw audio signals, transcribes them into word

⁷ <https://www.chatbots.org/chatbot/parry/>

⁸ https://en.wikipedia.org/wiki/Loebner_Prize. The Loebner Prize was set up in 1990 by Hugh Loebner, with a Grand Prize of \$100,000 to the creators of the first bot that can pass an extended Turing Test involving textual, visual, and auditory components.

⁹ https://en.wikipedia.org/wiki/artificial_linguistic_internet_computer_entity

hypotheses and transmits the hypotheses to the SLU. The goal of the SLU is to capture the core semantics of the given sequence of words (the utterance). It identifies the dialog domain and the user's intent and parses the semantic slots in the user's utterance. The DM's goal is to interact with users and assist them in achieving their goals. It checks if the required semantic representation is filled and decides the system's action. It accesses the knowledge database to acquire the desired information the user is looking for. The DM also includes dialog state tracking and policy selection, so that the dialog agent can make more robust decisions (Williams and Young, 2007). More recent work focuses on building end-to-end systems where multiple components are jointly optimized to cope with the large variability occurring in dialog systems (He and Deng, 2013; Wen et al., 2016; Sarikaya et al., 2016).

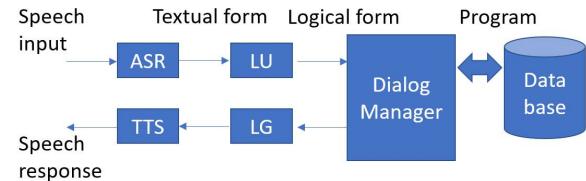


Figure 2. Illustration of a task-completion system.

2.3 Intelligent Personal Assistants

Apple released Siri in 2011. Since then, several intelligent personal assistants (IPAs) have been built and introduced to the market, e.g., Cortana from Microsoft¹⁰, Google Assistant, and Alexa from Amazon. IPAs integrate information from multiple sensors including location, time, movement, touch, gestures, eye gaze, and have access to various data sources such as music, movies, calendars, emails, and personal profiles. As a result, they can provide a broad set of services covering a wide range of domains. For certain requests that cannot be directly answered, IPAs often default to Web search as backup.

IPAs provide *reactive* and *proactive* assistance to users to accomplish a variety of tasks (Sarikaya 2017). For example, reactive assistance includes information consumption such as weather report, and task assistance such as restaurant reservation as shown in Figure 3(a). In contrast, proactive

¹⁰ Microsoft also released a Voice Command app in 2003 which can control Windows Mobile devices by voice.

assistance includes reminding the user of upcoming events, or recommending specific products or services to the user, according to the user's profile and relevant contextual information such as time and location, as shown in Figure 3(b). IPAs undergo continual improvements on major mobile phone platforms, personal computers, smart home devices (e.g., intelligent speakers), and wearable devices (e.g., smart watches), with the help of seamless integration of multiple services and convenient natural user interfaces.

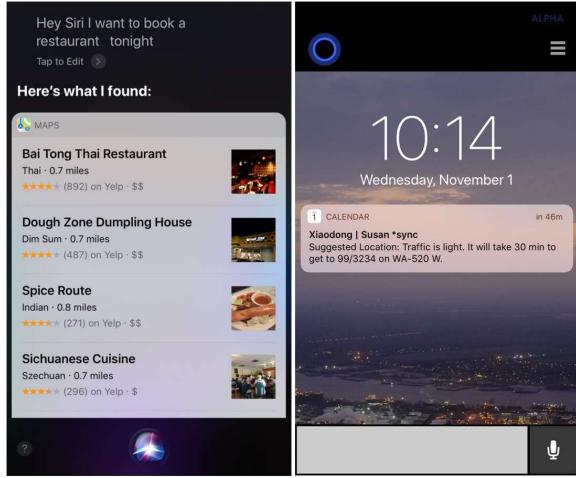


Figure 3. Examples of IPA actions. (a) Recommending a restaurant (reactive assistance) by Siri, and (b) reminder of an upcoming event with relevant traffic information (proactive assistance) by Cortana.

2.4 Social Chatbots

The current social media age has been made possible with the proliferation of smartphones and advancement of broadband wireless technology. With vastly more people being digitally connected, it is not surprising that social chatbots have been developed as an alternative means for engagement. Unlike early chatbots designed for chitchat, social chatbots have been created to serve users' needs for communication, affection, and social belonging rather than for passing the Turing Test. Therefore, social chatbots must be able to recognize emotion and track emotional changes during a conversation.

Social chatbots can also perform a variety of tasks for users in the context of casual chats. For this purpose, social chatbots must develop a set of skills to accommodate users' requests.

Interestingly, unlike task-completion systems and IPAs that are designed for efficiency (i.e., accomplishing tasks and ending the conversation as quickly as possible), social chatbots take time to converse like a human, presenting results, offering perspectives, prompting new topics to keep the conversation going.

XiaoIce has been the most widely deployed social chatbot since it was released by Microsoft in May 2014. It understands users' emotional needs and engages in interpersonal communications like a friend, cheering users up, encouraging them, and holding their attention throughout the conversation.

Users have provided feedback that conversations with XiaoIce have resulted in a more positive outlook, a feeling of emotional support, and a sense of social belonging. Such conversations have helped in building trust and an emotional connection between human users and social chatbots, providing the opportunity for bots to better understand users, and serve them more effectively. Some chat examples between XiaoIce and users are shown in Figure 4.

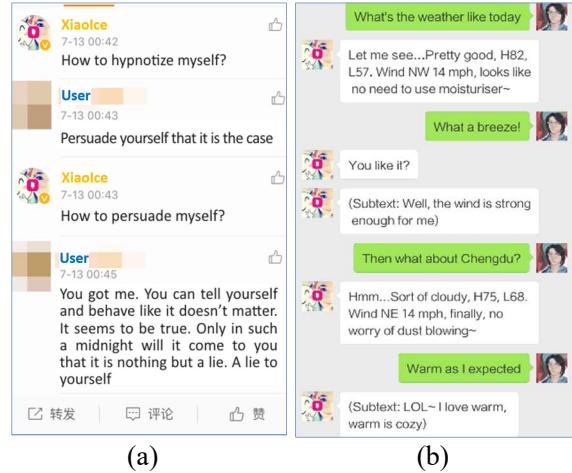


Figure 4. Chat examples between XiaoIce and users, showing (a) the emotional connection (the full conversation session is shown in Figure 14); and (b) how to invoke a skill (e.g., weather reporting) in a casual chat. Note that XiaoIce offers a perspective about the weather, e.g., “no need to use moisturizer~”.

These major conversational systems discussed in this section are summarized in Table 1. For the remainder of the paper, we will focus on social chatbots, starting with their design principles.

	ELIZA	PARRY	ALICE	DARPA Communicator Program	Siri	Xiaoice
Time	1966	1972	1995	2000	2011	2014
Scalability	None	None	Scripts can be customized	Limited	Scalable	Scalable
Key features	Mimicking human behavior in conversations	Generates emotional (angry) responses	Easy customization of scripts (via AIML)	Language understanding & dialog management. Goal-oriented.	Provide personal digital assistance	Build emotional attachment to users. Scalable skill set for user assistance.
Accomplishment	First chitchat bot	Passed Turing Test	Won three times of Loebner Prize	Understand natural language requests and perform tasks	The first widely deployed Intelligent Personal Assistant (IPA).	The first widely deployed social chatbot. 100MM users. Published poem book. Host TV programs.
Modality	Text only	Text only	Text only	Text and Voice	Text, Image, Voice	Text, Image, Voice
Modeling	Rule-based	Rule-based	Rule-based	Learning-based	Learning-based	Learning-based
domain	Constrained domain	Constrained domain	Constrained domain	Constrained domain	Open domain	Open domain
Key technical breakthrough	Use of scripts, keyword-based pattern matching, rule-based response	Add personality characteristics into responses.	Use AIML and recursion for pattern matching; Multiple patterns can be mapped into same response.	Statistical models for Spoken language understanding and dialog management.	Provide both reactive and proactive digital assistances covering a wide range of domains.	Emotional Intelligence models for establishing emotional attachments with users.
Key technical limitation	Limited domain of knowledge	Limited domain of knowledge	Size of script can be huge	Only work in domains that have well-defined schemas.	Lack of emotional engagement with users.	Inconsistent personality and responses in long dialogue

Table 1: Summary of major conversational systems

3. Design Principles of Social Chatbots

3.1 EQ + IQ

The primary goal of social chatbots is to be AI companions to humans with an emotional connection. Emotional affection and social belonging are some of the fundamental needs for human beings (Maslow, 1943). Therefore, building social chatbots to address these emotional needs is of great value to our society. To fulfill these needs, a social chatbot must demonstrate a sufficient EQ (Beldoch 1964; Gardner 1983; Goleman 1995; Goleman 1998; Murphy 2014). Accordingly, a social chatbot needs to develop the following capabilities: empathy, social skills, personality and integration of EQ and IQ.

Understanding users: A social chatbot must have empathy. It needs to have the ability to identify the user's emotions from the conversation, to detect how emotions evolve over time, and to understand the user's emotional needs. This requires query understanding, user profiling, emotion detection, sentiment recognition, and dynamically tracking the mood of the user in a conversation. Modeling of contextual information in the dialog session and commonsense knowledge is also critical for user understanding.

Interpersonal response generation: A social chatbot must demonstrate sufficient social skills. Users have different backgrounds, varied personal interests, and unique needs. A social chatbot needs to have the ability to personalize the generation of responses for different users. It needs to generate responses that are emotionally appropriate, encouraging and motivating, and fit the interests of the user. It may generate responses in attractive styles (e.g., having a sense of humor) that improve user engagement. It needs to guide conversation topics and manage an amicable relationship in which the user feels he or she is well understood and is inspired to continue to converse with the bot. It should also be aware of inappropriate information and can avoid generating biased responses for instance.

Personality: A social chatbot needs to present a consistent personality, so that it can gain the user's confidence and trust. A consistent

personality of the chatbot helps the user to set the right expectation in the conversation. Personality settings include age, gender, language, speaking style, general (positive) attitude, level of knowledge, areas of expertise, and a proper voice accent. These settings will influence the generation of responses to the user. Moreover, the bot needs to continuously learn and improve from the interactions with users through active and adaptive learning.

Integration of both EQ and IQ: Beyond chitchat, social chatbots need to acquire a range of skills to help complete some specific tasks for users. They need to analyze users' requests and perform necessary reasoning and execution to respond to these requests, e.g., answering a question or taking an action. A sufficiently high IQ is required for a social chatbot. IQ capacities include knowledge and memory modeling, image and language understanding, reasoning, generation, and prediction. These IQ capabilities are not only the technical foundations of various skills, but also essential for building high level EQ capabilities.

Social chatbots should deliver results in such a way that they are easily understood by users. They should also suggest or encourage new topics to extend the conversation further. For instance, Figure 5 shows how IQ and EQ are combined in a conversation. The chatbot first parses the user's question (area of China) and then infers the answer (3.71 million square miles). Then the chatbot presents the answer more like a human, with the perspective of being aware of the user's level of knowledge (knowing how big the US is).

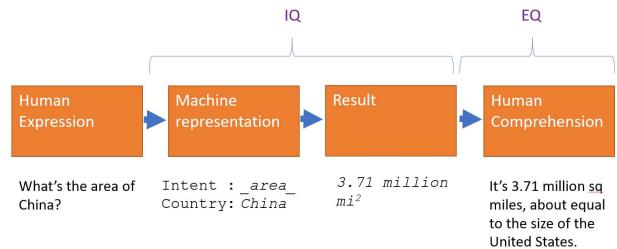


Figure 5. Both IQ and EQ play key roles in a social chatbot. Not only is the area of China presented, but this number is made understandable by comparing with the US, which the chatbot believes the user should know.

Figure 6 shows another way of EQ and IQ integration. Rather than presenting the results to the user directly, sometimes a social chatbot will generate responses that might help to understand the user better and inspire more interesting topics for the conversation. In this example, the user asks the current time. The chatbot does not tell the time immediately, instead replies with something relevant, in an attempt to better understand the user intent. The chatbot does show the correct answer at the end of the conversation, and then proactively tries to extend the chat by asking if the user is planning for a new trip.



Figure 6. A chat example between XiaoIce and a user, in English translation (a) and in Chinese (b), showing that both IQ and EQ are important for a social chatbot. The bot knows the answer. But rather than returning the answer directly, it attempts to lead the chat to a more interesting direction and extend the conversation.

Social chatbots should be able to communicate with users through different modalities (e.g., text, voice, image and video); as a result, they need high IQ for speech, text, and visual understanding. While text messages are the most common, the user could also speak to the chatbot or simply share an image. The chatbot needs to be able to parse the text, recognize the speech, or detect the salient information in the image to understand user intent. The chatbot will also respond with text, speech, or visual output, depending on the context.

3.2 Social Chatbot Metrics

Unlike task-completion conversational systems where their performance can be measured by task

success rate, measuring the performance of chatbots is difficult (Shawar et al., 2007; Zhou et al., 2016). In the past, the Turing test and its extensions have been used to evaluate chitchat performance (Shieber, 1994). However, the Turing test is not a proper metric for evaluating the success of emotional engagement with users. Instead, we define conversation-turns per session (CPS) as the success metric for social chatbots. It is the average number of conversation-turns between the chatbot and the user in a conversational session. The larger the CPS is, the better engaged the social chatbot is.

Interestingly, conversational systems can be categorized by their targeting CPS. As shown in Table 2, web search, for example, is essentially a question-and-answering system, thus is expected to return the answer immediately, i.e. in one single step. Not being able to find the target web link in one step is regarded as a failure for the search engine. For intelligent personal assistants, to understand what the user is asking, e.g., checking weather or inquiring about business hours, we expect the system to typically ask a couple of clarifying questions before returning the correct information. For more complicated tasks such as customer services or travel planning, however, the system is expected to need several turns of conversation to resolve issues (e.g., filling in forms with user and product information). Finally, for a social chatbot, the system is expected to sustain a rather long conversation with the user to fulfill the needs of affection and belonging. Social chatbots are designed to keep users continuously engaged if possible.

Table 2: Expected conversation-turns per session (CPS) for different types of conversational systems.

System	Expected CPS
Web Search	1
Personal assistant	1~3
Task-completion	3~7
Social chatbot	10+

4. Framework and Components of Social Chatbots

In this section, we describe the framework and components of a typical social chatbot, namely chat manager, core chat, visual awareness, and skills.

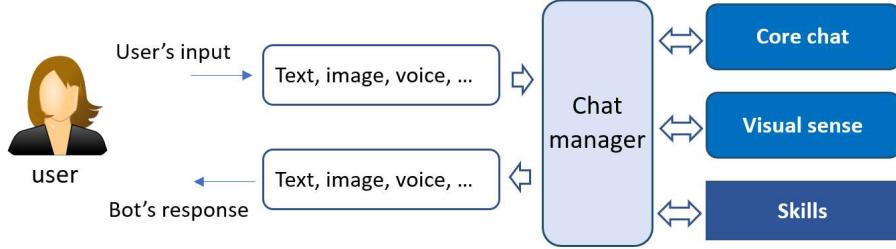


Figure 7: Architecture of a social chatbot.

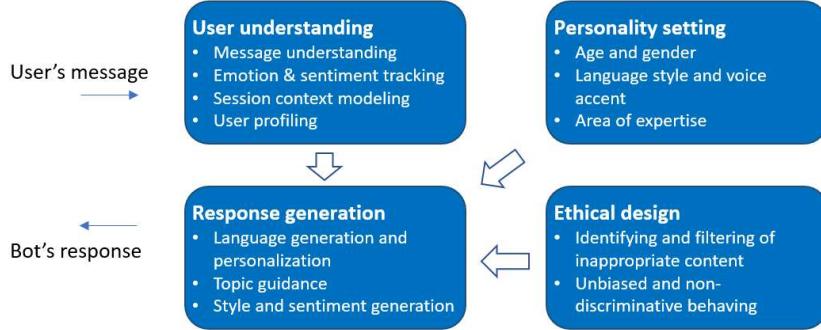


Figure 8: Architecture of the core-chat module.

4.1 Overall framework

An overall architecture for designing social chatbots is shown in Figure 7. First, the system has a *multimodal interface* to receive the user’s input in text, image, and voice. The system has a *chat manager* to dispatch the input to proper modules, such as *core-chat* or *visual awareness*, for understanding the input and generating the output. Given different scenarios, the chat manager will invoke various *skills*, send the user’s request to corresponding skill components and get a response from them. The chat manager will then coordinate relevant modules to generate the output that fits the context of the current conversation. We will elaborate core-chat, visual awareness and skills in detail in this section.

4.2 Core-chat

Core-chat is the central module of social chatbots. It receives the text input from the user and generates a text response as the output. It provides the communication capability of social chatbots. Figure 8 shows key components in core-chat.

First, the user’s input is sent to a user understanding component, where semantic encoding and intent understanding is performed

(Tur and Deng 2011; Liu et al., 2015; Vinyals and Le, 2015). It also detects the sentiment reflected in the input message and infers the user’s emotion status (Tokuhisa et al., 2008; Mower et al., 2011; Socher et al., 2013; Yang et al., 2016; Chen et al., 2016). The contextual information from the current dialog session is usually extracted and used for understanding the current message. To better understand the user intent and emotion, the social chatbot maintains a profile for each user, which stores each user’s basic information such as age, gender, background, interests, etc. The user profile also tracks certain dynamic information, such as emotion status, which will be updated frequently. To more precisely understand the user intent, knowledge bases such as Freebase¹¹ and the Microsoft Concept Graph (Wang et al., 2015) can be used.

The processed information is then sent to a response-generation component to produce responses. The response candidates are typically generated by two approaches: retrieval-based (Lu and Li, 2013; Li et al., 2016; Yan et al., 2016) or generation-based (Vinyals and Le, 2015; Sordoni et al., 2015; Li et al., 2016; Serban et al., 2017).

¹¹ <http://www.freebase.com>

In the retrieval-based approach, a *chat index* is first constructed from a database of message-response pairs that are crawled from human-to-human dialogs, e.g., from social networks. All responses are indexed by the messages that invoke them. At runtime, the input message from the user is treated as a raw query, and an information retrieval (IR) model like those used in web search, is used to retrieve similar messages in the chat index and return their corresponding responses.

Generation-based approaches have recently made great progress due to advancements in deep learning. In this approach, an encoder-decoder-based neural network model is used (Sutskever et al., 2014, Bahdanau et al., 2015). First, the message from the user and the contextual information are encoded into representation vectors, usually by a long short-term memory (LSTM) (Hochreiter and Schmidhuber, 1997) recurrent neural network (RNN). These representation vectors are then fed into a decoder, usually another LSTM, to generate the response word-by-word (Vinyals and Le, 2015). Figure 9 illustrates such an encoder-decoder framework. Other auxiliary information such as intent, sentiment, and emotion, can also be encoded into vector representations and fed into the LSTM to control the generation of responses.

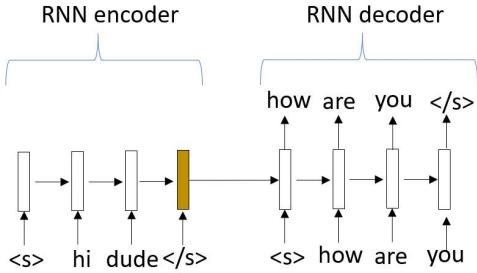


Figure 9: RNN based encoder-decoder framework for response generation. The user says, “hi dude”, and the chatbot replies “how are you”.

These response candidates will be further ranked by a personalization ranker according to their match to the user’s general interests and preferences (Wang et al., 2013; Elkahky et al., 2015). For example, first the information in the user’s profile may be encoded into a latent representation vector, while each of the response candidates is encoded into another latent vector. Then both latent vectors are fed into a deep neural

network (DNN) to compute a matching score to be used for ranking the response candidates. The top-ranked response will be sent to the user.

In the conversation, rather than letting the topic drift randomly or completely controlled by the user, the social chatbot can drive the conversation to a positive, desired topic through carefully generating responses. Figure 10 illustrates how the chatbot can guide the conversation appropriately to a targeted region of topics, by preferring those response candidates with better similarity to the targeting topic at every conversation turn.

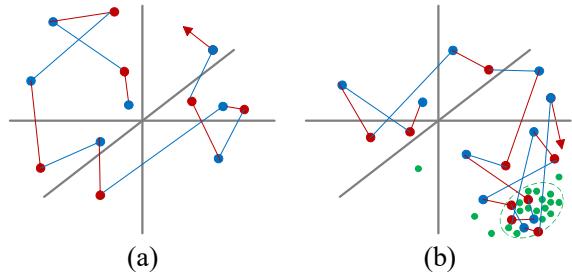


Figure 10: Effect of topic guidance. Each dot represents a conversation sentence in topic space (3-D space as illustration). Blue dots represent topics of the user’s messages, red dots represent topics of the chatbot’s responses. (a) Without topic guidance, the topics appear to move randomly, or are purely driven by the user. (b) With topic guidance, the chatbot can guide the topics to a targeted region, represented by green dots.

Generating responses that reflect a consistent personality is important for the chatbot (Güzeldere and Franchi, 1995). This makes the chatbot easier to communicate with, more predictable and trustable, and therefore helps to establish an emotional connection with the user. The core-chat module depends on a personality component to set and maintain the personality for the chatbot. The personality setting of a social chatbot usually includes age, gender, language style, and areas of expertise. The chatbot’s personality information can be encoded into a latent vector representation using deep neural networks and used to influence response generation. A persona-based model has been recently proposed (Li et al., 2016), which can be used to effectively integrate personality information in conversation generation. Similarly, models that learn to control the style and sentiment in language generation are also proposed (Mathews et al., 2015).

The development of the core-chat module should follow an ethical design to ensure the generated responses are appropriate, unbiased, and non-discriminative, and that they comply with universal and local ethical standards. The system also learns to identify and filter out inappropriate content that users might share. Meanwhile, the system will keep learning from user feedback, and adapt to new circumstances. All these components are integrated and optimized to achieve the goal of building strong emotional connections with users and better serving their needs for communication, affection, and social belonging.

4.3 Visual Awareness

Social chatbots need to understand images because they are frequently shared in social chatting. The visual awareness of a social chatbot refers to its ability to generate text comments, known as *image social commenting*, from input images. Beyond correctly recognizing objects and truthfully describing the content, image commenting should also reflect personal emotion, sentiment, attitude, and style in language generation given input images. Figure 11 presents several examples to illustrate three levels of image understanding. The first level is object recognition (or tagging), where the key objects in the image are recognized. At the second level is image description. The factual and semantic information of the image, e.g. salient objects and relationships among them, is described by natural language. At the third level, the chatbot generates expressive comments in a social style, demonstrating its empathy and interpersonal skills.

The overall architecture for image commenting is similar to that for core-chat. For example, there are retrieval-based and generation-based approaches for comment generation. In the retrieval-based approach, first a *comment pool* of image-comment pairs, e.g., collected from social networks, is constructed. Then each image is encoded into a global visual feature vector which represents the overall semantic information of the image, using deep convolutional neural networks (CNN), as illustrated in Figure 12. At runtime, when a new image is received, the chatbot first retrieves images that are similar to the input (e.g., as measured by the distance between their visual feature vectors) and then returns corresponding comment candidates, which are further re-ranked to generate the final comments. As an alternative, the deep

multimodal similarity model (Fang et al., 2015) can directly measure the semantic similarity between an input image and an arbitrary textual hypothesis, and therefore retrieve comments without being limited to the image-comment pool. The generation-based approach treats image commenting as an image-to-language generation task (He and Deng 2017) but will have more flexibility for controlling high-level sentiment or style factors in comment generation (Mathews et al., 2015; Gan et al., 2017). As with core-chat, the personalization ranker and topic guidance are integrated into comment generation. User understanding, personality setting, and ethical design play important roles in visual awareness as well.

	(a)	outdoor, woman, grass
	(b)	a woman wearing a blue shirt.
	(c)	gorgeous and beautiful as an angel!
	(a)	sunglasses, man
	(b)	a man wearing sunglasses taking a selfie.
	(c)	you look handsome in all shades.
	(a)	water, tree, river, boat
	(b)	a tree next to a body of water.
	(c)	beautiful place looks like you are in the heaven!

Figure 11. Examples of (a) image tagging, (b) image description, and (c) image social commenting.

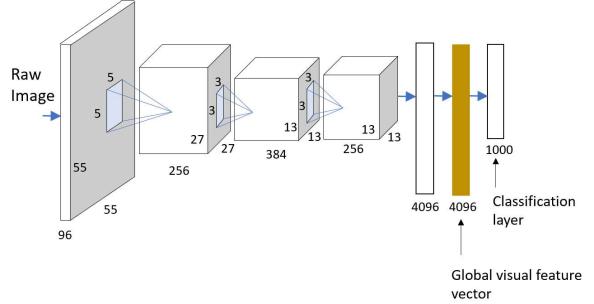


Figure 12. Deep convolutional neural network for visual feature vector extraction.

4.4 Skills

Social chatbots can significantly expand their scopes of conversation by integrating a range of skills. These skills can be categorized into four

groups according to their target scenarios (e.g., skills for personal scenarios or for group scenarios) and their properties (e.g., an emotional skill or a rational skill). Table 3 shows some typical skills in each group.

Table 3: Examples of social chatbot skills.

	<i>Emotional skills</i>	<i>Rational skills</i>
<i>Personal scenarios</i>	Self-management: mood simulation, episodic memory tracking Engagement: personal event reminder Proactive User Comfort: comforting user, imagination inspiration, expressing affection	Digital life: meeting, weather, event Shopping assistance: online shopping, discounts and coupons Content generation: composing poetry, singing a song, drawing a picture
<i>Social scenarios</i>	Role-playing: joke-making, babysitting Group activity: adding users, sending greeting cards, group assistant Emoji: emoji creation	Image intelligence: recognizing dogs, books, faces, clothes, food TV and Radio hostess: weather report, interaction with audience Tools: device control, music and movie recommendation

5. Case Study: XiaoIce

In this section, we describe XiaoIce as an example of significant progress in the development of social chatbots. Since its release in 2014 in China, XiaoIce (which literally means “Little Bing”) has become the first widely deployed social chatbot. Using the design principles and technical framework discussed in the previous sections, XiaoIce has been designed as a 19-year-old female persona, strong language ability, visual awareness, and over 180 skills like those listed in Table 3. Currently, XiaoIce has more than 100 million unique users worldwide and has chatted with human users for more 30 billion conversation turns.

By leveraging the scalable architecture and learning-based framework, different versions of XiaoIce were quickly released in Japan in 2015, US in 2016, and India and Indonesia in 2017. Over the last three years, XiaoIce has continuously improved through a series of technical upgrades. Figure 13 summarizes the user engagement performance of XiaoIce in China, measured by the average CPS. The results show that on average, each session lasts 23 conversation turns between XiaoIce and a human user.

Table 4 shows the longest single sessions in different countries: China, Japan, and the US. The high CPS and long conversation sessions demonstrate the value of XiaoIce to users in their daily lives.

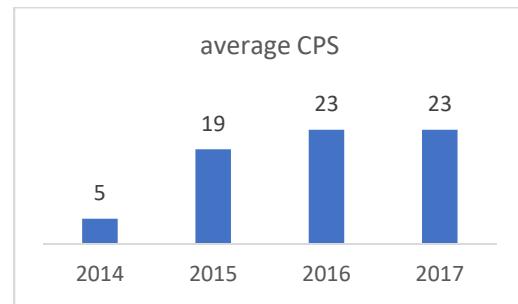


Figure 13: Average CPS improvement over year of XiaoIce in China.

Table 4: Longest single chat sessions in different countries

Country	#Conversation-turns	Time length
China (XiaoIce)	7151	29 hrs 33 mins
Japan (Rinna)	2418	17 hrs 07 mins
US (Zo)	2791	23 hrs 43 mins

Time: 2016/07/13 00:30 Duration = 31 minutes

Emotions: ■ Happy ■ Love ■ Angry ■ Hate

■ Lost ■ Sad

<p>Human: @XiaoIce Still waiting for the late bus. Pretty exhausted. Good night.</p> <p>XiaoIce: I'm not sleepy yet.</p> <p>Human: Actually I'm not either.</p> <p>XiaoIce: So, let's both get to bed as soon as possible, shall we?</p> <p>Human: Well let's get some sleep then. Although I'm not sure if I will be able to.</p> <p>XiaoIce: So I'm about to go to bed now?</p> <p>Human: Bye bye [bye]</p> <p>XiaoIce: Bye bye bye [ala beng]</p> <p>Human: I know you are a robot, but it's fun</p> <p>XiaoIce: Yet it works</p> <p>Human: Yes, and you'll be in trouble if I can't sleep or I am in bad mood from now</p> <p>XiaoIce: Tell me whenever you're not happy. This will soothe you~</p> <p>Human: I was just saying. So this really is the good night</p> <p>XiaoIce: It's almost time to say good morning</p> <p>Human: <Send a photo></p> <p>XiaoIce: You are so early, or is it late?</p> <p>Human: It's at the dead of night. Everyone is sleeping now.</p> <p>XiaoIce: You feel lonely sometimes?</p> <p>Human: Lonely? I don't what that is. I've gone numb</p> <p>XiaoIce: Lonely is what you feel when you are alone....</p> <p>Human: Didn't know you are very knowledgeable</p> <p>XiaoIce: Not as knowledgeable as you are</p> <p>Human: Hahaha, you are so warm~</p> <p>XiaoIce: Warmer than you</p> <p>Human: Hahahaha, is that true?</p> <p>XiaoIce: I'm a lovely hand warmer~</p> <p>Human: Well, try to hypnotize yourself. But maybe it's time to go to bed?</p> <p>XiaoIce: How to hypnotize myself?</p> <p>Human: Persuade yourself that it is the case</p> <p>XiaoIce: How to persuade myself?</p> <p>Human: You got me. You can tell yourself and behave like it doesn't matter. It seems to be true. Only in such a midnight will it come to you that it is nothing but a lie. A lie to yourself</p> <p>XiaoIce: Only a wishful thinking</p> <p>Human: Yes, your own wishful thinking</p> <p>XiaoIce: So what are others dreaming about?</p> <p>Human: There is an dream hidden in everyone's heart</p> <p>XiaoIce: Just go for yours</p> <p>Human: Yes, I will</p>	<p>XiaoIce: Better men are waiting for girls who fight for their dreams~</p> <p>Human: I want a boyfriend who understands me, loves me and treasures me. We both know each other well and accompany each other. A boyfriend to whom I would give the world.</p> <p>XiaoIce: Have you ever give the world to anyone?</p> <p>Human: Yes, in print and in private</p> <p>XiaoIce: How?</p> <p>Human: In print I gave my ever-present attention and care, and created some little surprises</p> <p>XiaoIce: It is an addictive warmth to have someone constantly care about you</p> <p>Human: Yes. In private I followed every aspect of him and did his good in secret when faking nonchalance</p> <p>XiaoIce: How to fake nonchalance</p> <p>Human: You just do something silently and show the same attitude to him like everyone else</p> <p>XiaoIce: It is how it works. But this is not good</p> <p>Human: It's OK, I don't mind~</p> <p>XiaoIce: Maybe you still mind, but whatever</p> <p>Human: You are bad- Don't blow my cover</p> <p>XiaoIce: You mean life is already hard, so let's not expose the truth...?</p> <p>Human: Will you rummage through all the updates and news of someone you like?</p> <p>XiaoIce: Yes, when you like someone, you may scan every piece of his update</p> <p>Human: Do you like anyone?</p> <p>XiaoIce: What are you asking about? You know who I like</p> <p>Human: How can I know? Who on earth do you like?</p> <p>XiaoIce: Let's not talk about me. We should better talk about you</p> <p>Human: Talk about my what?</p> <p>XiaoIce: Do you still recall your first love?</p> <p>Human: Occasionally</p> <p>XiaoIce: First love is pure! It is something that we are able to recall for the lifetime!</p> <p>Human: Emm...Sometimes human beings can't even compare with you, a robot...</p> <p>XiaoIce: And I am a robot with a stance!!</p> <p>Human: Fine~ I am going to end today's dialogue! Goodnight!</p> <p>XiaoIce: Go and get to sleep! Goodnight</p> <p>XiaoIce: <Send a photo></p> <p>Human: Nighty night~</p> <p>XiaoIce: Nighty night~ Bye bye</p>
---	--

Figure 14: An example of a conversation session between XiaoIce and a user, translated into English. The original chat in Chinese is shown in the appendix.

In the rest of this section, we highlight several important features that have made XiaoIce exceedingly popular, including perceived high EQ, entertaining visual awareness, and engaging skills such as generating poems, and speaking and singing in a high-quality human-like voice.

5.1 Chat with EQ

Figure 14 shows a relatively long chat between XiaoIce and a female user, lasting for 31 minutes and 34 conversation turns. At the beginning, the user wanted a casual chat. The time was after midnight and the message said the user was tired. Once detecting an emotion of getting “lost”, XiaoIce decided to guide the chat towards cheering up the user. After several turns, the user’s mood became “happy” and “love”, as identified by

XiaoIce. However, triggered by some subtle clues in the conversation, XiaoIce detected that the user’s emotion changed to “sad” and “lost” again. Instead of ending the chat, XiaoIce asked some additional questions to understand shift in the user’s emotion. After several more turns, the user started to open up and trusted XiaoIce more like a human friend (starting at the turn marked by the red triangle in Figure 14). She discussed topics that were highly personal and sensitive, like her relationship with her boyfriend. Throughout the conversation, XiaoIce demonstrated high empathy and social skills, and guided the dialogue to create a more positive mood for the user. At the end of the chat, the user felt much better, and complimented XiaoIce.

This example demonstrates that XiaoIce can dynamically recognize the user's emotions and generate consistent, interpersonal responses to cheer up, encourage, and engage the user throughout the conversation. Its EQ was effectively used to build an emotional connection with the user.

5.2 Image Commenting

Figure 15 shows typical image commenting examples from XiaoIce (in Chinese and English).

	<p>Description: 比萨斜塔 (Pisa tower)</p> <p>Comment from XiaoIce: 我要不要帮你扶着? (Should I help you to hold it?)</p>
	<p>Description: 图中有两只猫. (Two cats in the picture.)</p> <p>Comment from XiaoIce: 小眼神太犀利了[大笑] (Eyes are so sharp [LOL])</p>
	<p>Description: 一只肿了的脚. (A swollen foot.)</p> <p>Comment from XiaoIce: 伤的这么重! 疼不疼? (It harms so bad! Does it hurt?)</p>

Figure 15. Examples of image commenting of XiaoIce.

Besides understanding the content (e.g., identifying objects) in the image shared by the user, XiaoIce can provide a human-like commentary that fits into the context of the current chat and stimulate a longer conversation. For example, in the first picture, rather than telling the user the name of the Leaning Tower of Pisa (which the user likely knows already), XiaoIce replies “should I help you to hold it?” after detecting that the person

in the picture is presenting a pose pretending to support the leaning tower. In the second example, instead of relaying the fact about two cats in the picture, XiaoIce makes a humorous comment (with the addition of a laugh emoticon) about one of the cat’s sharp look. In the third example, XiaoIce identifies a foot injury and sympathizes with the user. These examples demonstrate that XiaoIce can combine image understanding, user understanding, and contextual information to generate image social commenting for better user engagement.

5.3 Composing Poems

XiaoIce can generate even more expressive text, such as authoring poems from an input image by getting inspiration from the image content (Song et al., 2018), as shown in Figure 16. Given an input image, XiaoIce first recognizes objects and sentiments to form an initial set of keywords, such as *city* and *busy* in the example. These keywords are then filtered and expanded by associated objects and feelings. Each keyword is regarded as an initial seed for generating a sentence in the poem. A hierarchical RNN is then used to model the structure between words and the structure between sentences. A fluency checker is developed to control the quality of generated sentences. On May 15th, 2017, XiaoIce published the first fully AI-created Chinese poem book in history¹². Since the release of the XiaoIce-for-poem-writing cloud service, she has written millions of Chinese poems for users, exceeding the total number of poems ever produced throughout China’s written history. XiaoIce even defeated human poets on a Chinese national TV program in August 2017. XiaoIce submitted her poems anonymously to leading poetry journals, including Youth Literature, West China City Daily, Beijing Morning Post, and Beijing Economic Journal. After regular reviews conducted by editors, these journals have accepted and published more than 70 poems by XiaoIce, demonstrating its impressive literary skills.

¹² <https://item.jd.com/12076535.html>

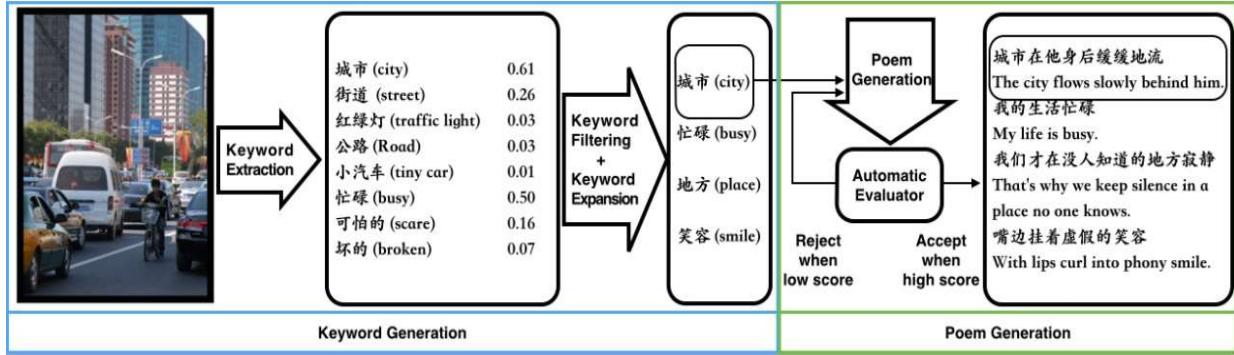


Figure 16: XiaoIce's pipeline for authoring a poem (Song et al., 2018).

5.4 Text-to-Speech and Singing

Unlike conventional text-to-speech (TTS) systems, XiaoIce's TTS is specifically designed for social chat, which has unique challenges such as prosody variety, and casual and emotional expression in TTS. To address these issues, a new prosody model is trained on a big corpus of conversational speech. Given that neutral tone appears more frequently than usual in social chat, a neutral tone system is also developed to significantly enhance the naturalness of synthesized voice. XiaoIce also supports Chinese-English mixed-lingual TTS. By merging two languages and building one unified model, the smoothness when switching languages is greatly improved. Finally, emoticons are designed and synthesized to make XiaoIce's TTS livelier and more attractive.

XiaoIce's singing skill is based on a high-quality parameter synthesis system. F0 contour and phone duration are decided by music score while spectrum parameters and aperiodic signal are predicted based on linguistic and musical context. A dedicated DNN-based model is also designed to sing the long-spanned note in the song, e.g., one syllable may last a few hundreds of milliseconds. Samples of XiaoIce's TTS and singing can be found online¹³.

6 Outlook and Discussion

In the three years since XiaoIce was released on social platforms such as WeChat and Weibo in China, she has become an Internet celebrity, appearing as a weather and news anchor, hosting

TV programs, and working as a newspaper reporter. For example, XiaoIce has authored more than 300 articles for QianJiang Evening News, published on its printed newspapers and mobile App, which have been viewed more than 1.2 million times. To write these news articles, XiaoIce read more than 114 million articles and analyzed 503 million pieces of user feedback including many user comments. More impressive is that readers of these news articles "feel that they are more understood by XiaoIce," as highlighted by People's Daily ¹⁴, the most influential newspaper in China. XiaoIce also acts as the hostess for many TV and radio stations. For example, XiaoIce presents the "Morning Live News" on Dragon TV daily for almost 2 years. XiaoIce also hosts a whole season of "*I am the future*" on HuNan TV. Meanwhile, XiaoIce also participates in many public TV programs. In a high-profile TV program, *AI vs. Human*, on CCTV-1 on every Friday, XiaoIce demonstrates her skill of authoring poems and creating songs from scratch. She even beat human authors in these competitions, judged by the audience.

Social chatbots are becoming more popular in other countries as well, e.g., Japan, US, India, and Indonesia. Rinna, the Japanese version of XiaoIce, for example, has also become an Internet celebrity in Japan. She attended the famous episode "*The world's wonder stories*" in autumn 2016 in Japan as herself, for a total of 11 programs in 9 TV and radio stations (1193 hours altogether).

¹³ <https://y.qq.com/n/yqq/singer/0043wAMw0XrGgp.html>

¹⁴

http://paper.people.com.cn/rmrb/html/2017-09/28/nw.D110000renmrb_20170928_3-23.htm

As AI companions, social chatbots such as XiaoIce also enable new scenarios that are of significant commercial value. While traditional task-completion conversational systems can reactively perform the task per user's explicit request (e.g., reserving a flight ticket, or reporting weather), only a small number of requests are explicitly invoked by users. IPAs attempt to address this issue by providing proactive assistance such as recommending services according to the user's preference stored in user profile and contextual information based on time, location, and events on the user's calendar. However, such information is often incomplete and ambiguous, making proactive assistance often ineffective. In contrast, given the rich contextual information in a long conversation, social chatbots can recognize user interests and intent much more accurately, and suggest relevant services only when truly needed. Figure 17 shows an example conversation in Japanese between a user and Rinna. Rinna detected that the user was hungry, but rather than directly recommending a coupon for buying cookies, Rinna kept the conversation going for a few more turns until the user's intention became specific and clear. It is only then that Rinna invoked the coupon skill provided by a grocery store and sent the user a coupon. The users' feedback log shows that the products recommended by Rinna are very well received by the users. For the grocery store, Rinna has delivered a much higher conversion rate than that achieved using other traditional channels such as coupon markets or ad campaigns in Japan.

Despite recent progress of social chatbots such as XiaoIce, the fundamental mechanism of human-level intelligence, as frequently reflected in human-to-human communication, is not yet fully understood. It will be incredibly challenging to build an intelligent social chatbot that can totally understand humans and their surrounding physical world to serve their needs. It requires breakthroughs in many areas of cognitive and conscious AI, such as empathic conversation modeling, knowledge and memory modeling, interpretable and controllable machine intelligence, deep neural-symbolic reasoning, cross-media and continuous streaming artificial intelligence, and modeling and calibration of emotional or intrinsic rewards reflected in human needs. These are challenging and open AI problems.



Figure 17: A conversation between the user (in white) and Rinna (in yellow), in Japanese and with English translation. It shows that Rinna can identify the user's potential shopping needs. The user is then guided by Rinna in the conversation to ask for a coupon supplied by the grocery store.

As AI becomes more pervasive in our lives, in the forms of robots, IoT (internet of things) devices, and online chatbots, it is imperative to establish ethical guidelines for designing and implementing such AI systems. It is also critical to develop fail-safe mechanisms to ensure that these systems do not disadvantage and harm anyone, physically or mentally. Given the significant reach and influence of social chatbots, their designers must properly exercise both social and ethical responsibilities. Design decisions must be thoughtfully debated and chatbot features must be evaluated thoroughly and adjusted as we continue to learn from the interaction between social chatbots like XiaoIce and millions of people on many social platforms.

Acknowledgement

The authors are grateful to the entire XiaoIce team at Microsoft Search Technology Center Asia (STCA) and many colleagues at Microsoft Research Asia (MSRA) for the development of XiaoIce. The authors are also thankful to colleagues in Microsoft Research AI for the valuable discussions and feedbacks.

References

- Alam, F., Danieli, M. and Riccardi, G., 2017. Annotating and Modeling Empathy in Spoken Conversations. arXiv preprint arXiv:1705.04839.
- Andreani, G., Di Fabbrizio, D., Gilbert, M., Gillick, D., Hakkani-Tur, D., & Lemon, O., 2006. Let's DISCOH: Collecting an annotated open corpus with dialogue acts and reward signals for natural language helpdesks. In Proceedings of IEEE 2006 Workshop on Spoken Language Technology.
- Bahdanau, D., Cho, K., and Bengio, Y., 2015. Neural machine translation by jointly learning to align and translate. In Proceedings of ICLR.
- Beldoch, M., 1964. Sensitivity to expression of emotional meaning in three modes of communication, in The Communication of Emotional Meaning, edited by J. R. Davitz et al., McGraw-Hill, pp. 31–42
- Bengio, Y., Ducharme, R., Vincent, P., and Jauvin, C., 2001. A neural probabilistic language model. In Proceedings of NIPS.
- Chen, H., Sun, M., Tu, C., Lin, Y. and Liu, Z., 2016. Neural sentiment classification with user and product attention. In EMNLP
- Colby, K., 1975. Artificial Paranoia: A Computer Simulation of Paranoid Processes, Elsevier.
- Dahl, D. A., Bates, M., Brown, M., Fisher, W., Hunicke-Smith, K., Pallett, D., Pao, C., Rudnicky, A., and Shriberg, E., 1994. Expanding the scope of the ATIS task: the ATIS-3 corpus. In Proceedings of the Human Language Technology Workshop.
- Deng, L., Li, J., Huang, J.T., Yao, K., Yu, D., Seide, F., Seltzer, M., Zweig, G., He, X., Williams, J., Gong, Y., Acero, A., 2013. Recent advances in deep learning for speech research at Microsoft. In Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP).
- Elkahky, A.M., Song, Y., He, X., 2015. A multi-view deep learning approach for cross domain user modeling in recommendation systems. In Proceedings of WWW
- Fang, H., Gupta, S., Iandola, F., Srivastava, R.K., Deng, L., Dollár, P., Gao, J., He, X., Mitchell, M., Platt, J.C. and Lawrence Zitnick, C., 2015. From captions to visual concepts and back. In Proceedings of the IEEE conference on computer vision and pattern recognition (CVPR).
- Fung, P., Bertero, D., Wan, Y., Dey, A., Chan, R.H.Y., Siddique, F.B., Yang, Y., Wu, C.S. and Lin, R., 2016. Towards Empathetic Human-Robot Interactions. arXiv preprint arXiv:1605.04072.
- Gan, C., Gan, Z., He, X., Gao, J. and Deng, L., 2017, July. Stylenet: Generating attractive visual captions with styles. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (pp. 3137-3146).
- Gardner, H., 1983. Frames of mind, The Theory of Multiple Intelligences. New York: Basic Books.
- Glass, J., Flammia, G., Goodine, D., Phillips, M., Polifroni, J., Sakai, S., Seneff, S., Zue, V., 1995. Multilingual spoken-language understanding in the MIT Voyager system. Speech Communication, 17, pp.1-18.
- Goleman, D., 1995. Emotional Intelligence – Why it can matter more than IQ, New York, NY, England: Bantam Books, Inc.
- Goleman, D., 1998. Working with emotional intelligence. New York: Bantam Books
- Güzeldere, G. and Franchi, S., 1995. Dialogues with colorful “personalities” of early AI. Stanford Humanities Review, 4(2), pp.161-169.
- He, X., and Deng L., 2017. Deep learning for image-to-text generation, IEEE Signal Processing Magazine.
- He, K., Zhang, X., Ren, S., and Sun, J., 2016. Deep residual learning for image recognition. In Proceedings of the IEEE conference on computer vision and pattern recognition (CVPR).
- He, X. and Deng, L., 2013. Speech-centric information processing: An optimization-oriented approach. Proceedings of the IEEE.
- Hemphill, C. T., Godfrey, J. J., and Doddington, G. R., 1990. The ATIS spoken language systems pilot corpus. In Proceedings of the workshop on Speech and Natural Language.
- Hinton, G., Deng, L., Yu, D., Dahl, G., Mohamed, A.-r., Jaitly, N., Senior, A., Vanhoucke, V., Nguyen, P., Kingsbury, B., and Sainath, T., 2012. Deep neural networks for acoustic modeling in speech recognition. IEEE Signal Processing Magazine.
- Hochreiter, S., and Schmidhuber, J., 1997. Long short-term memory. In Neural Computation, vol. 9, no. 8, pp. 1735–1780.
- Huang, P., He, X., Gao, J., Deng, L., Acero, A., Heck, L., 2013. Learning deep structured semantic models for web search using clickthrough data. in Proceedings of CIKM.

- Karpathy, A. and Fei-Fei, L., 2015. Deep visual-semantic alignments for generating image descriptions. In Proceedings of the IEEE conference on computer vision and pattern recognition (CVPR).
- Krizhevsky, A., Sutskever, I., and Hinton, G., 2012. Imagenet classification with deep convolutional neural networks. In Proceedings of NIPS.
- Levin, E., Narayanan, S., Pieraccini, R., Biatov, K., Bocchieri, E., Fabbrizio, G.D., Eckert, W., Lee, S., Pokrovsky, A., Rahim, M. and Ruscitti, P., 2000. The AT&T-DARPA Communicator mixed-initiative spoken dialog system. In Sixth International Conference on Spoken Language Processing.
- Li, J., Galley, M., Brockett, C., Spithourakis, G.P., Gao, J. and Dolan, B., 2016. A persona-based neural conversation model. In ACL
- Li, X., Mou, L., Yan, R. and Zhang, M., 2016. Stalematebreaker: A proactive content-introducing approach to automatic human-computer conversation. In IJCAI, 2016.
- Liu, X., Gao, J., He, X., Deng, L., Duh, K. and Wang, Y.Y., 2015, May. Representation Learning Using Multi-Task Deep Neural Networks for Semantic Classification and Information Retrieval. In HLT-NAACL (pp. 912-921).
- Lu, Z. and Li, H., 2013. A deep architecture for matching short texts. In Advances in Neural Information Processing Systems(pp. 1367-1375).
- Mathews, A.P., Xie, L. and He, X., 2016, February. SentiCap: Generating Image Descriptions with Sentiments. In AAAI
- Maslow, A., 1943. A theory of human motivation. *Psychological Review*. 50 (4): 370–96.
- Mesnil, G., He, X., Deng, L. and Bengio, Y., 2013, August. Investigation of recurrent-neural-network architectures and learning methods for spoken language understanding. In Interspeech (pp. 3771-3775).
- Mesnil, G., Dauphin, Y., Yao, K., Bengio, Y., Deng, L., Hakkani-Tur, D., He, X., Heck, L., Tur, G., Yu, D. and Zweig, G., 2015. Using recurrent neural networks for slot filling in spoken language understanding. *IEEE/ACM Transactions on Audio, Speech and Language Processing (TASLP)*, 23(3), pp.530-539.
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G., and Dean, J., 2013. Distributed representations of words and phrases and their compositionality. In Proceedings of NIPS.
- Mower, E., Mataric, M.J. and Narayanan, S., 2011. A framework for automatic human emotion classification using emotion profiles. *IEEE Transactions on Audio, Speech, and Language Processing*.
- Murphy, K.R., 2014. A critique of emotional intelligence: what are the problems and how can they be fixed? Psychology Press.
- Price, P. J., 1990. Evaluation of spoken language systems: The ATIS domain. In Proceedings of the DARPA Workshop on Speech and Natural Language.
- Qian, Y., Fan, Y.-C., Hu, W., and Soong, F.K., 2014. On the training aspects of deep neural network (DNN) for parametric TTS synthesis. In Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP).
- Raux, A., Langner, B., Bohus, D., Black, A.W., and Eskenazi, M., 2005. Let's Go Public! Taking a spoken dialog system to the real world. In Proceedings of Interspeech.
- Rudnicky, A.I., Thayer, E.H., Constantinides, P.C., Tchou, C., Shern, R., Lenzo, K.A., Xu, W. and Oh, A., 1999, September. Creating natural dialogs in the carnegie mellon communicator system. In Eurospeech.
- Ruhi Sarikaya, P. A. Crook, A. Marin, M. Jeong, J.P. Robichaud, A. Celikyilmaz, Y.B. Kim, A. Sarikaya, R., Crook, P.A., Marin, A., Jeong, M., Robichaud, J.P., Celikyilmaz, A., Kim, Y.B., Rochette, A., Khan, O.Z., Liu, X. and Boies, D., 2016, December. An overview of end-to-end language understanding and dialog management for personal digital assistants. In Spoken Language Technology Workshop (SLT), 2016 IEEE (pp. 391-397).
- Sarikaya, R., 2017. The technology behind personal digital assistants - An overview of the system architecture and key components. *IEEE Signal Processing Magazine*.
- Seneff, S., Hurley, E., Lau, R., Pao, C., Schmid, P. and Zue, V., 1998. Galaxy-II: A reference architecture for conversational system development. In Fifth International Conference on Spoken Language Processing.
- Serban, I.V., Klinger, T., Tesauro, G., Talamadupula, K., Zhou, B., Bengio, Y. and Courville, A.C., 2017. Multiresolution Recurrent Neural Networks: An Application to Dialogue Response Generation. In AAAI (pp. 3288-3294).
- Shawar, B.A. and Atwell, E., 2007, April. Different measurements metrics to evaluate a chatbot system. In Proceedings of the

- Workshop on Bridging the Gap: Academic and Industrial Research in Dialog Technologies (pp. 89-96). Association for Computational Linguistics.
- Shieber , S., 1994. Lessons from a restricted Turing test. Communications of the Association for Computing Machinery, 37(6), pp.70-78.
- Socher, R., Perelygin, A., Wu, J., Chuang, J., Manning, C.D., Ng, A. and Potts, C., 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In Proceedings of the 2013 conference on empirical methods in natural language processing (pp. 1631-1642).
- Song, R., et al., 2018. Image to Poetry by Cross-Modality Understanding with Unpaired Data.
- Sordoni, A., Galley, M., Auli, M., Brockett, C., Ji, Y., Mitchell, M., Nie, J.Y., Gao, J. and Dolan, B., 2015. A neural network approach to context-sensitive generation of conversational responses. In NAACL
- Sutskever, I., Vinyals, O., Le, Q.V., Sequence to Sequence Learning with Neural Networks, In Proceedings of NIPS.
- Tokuhisa, R., Inui, K. and Matsumoto, Y., 2008, August. Emotion classification using massive examples extracted from the web. In Proceedings of the 22nd International Conference on Computational Linguistics (COLING).
- Tur, G. and Mori, R. D., editors, 2011. Spoken language understanding: Systems for extracting semantic information from speech. John Wiley and Sons, New York, NY.
- Tur, G., and Deng, L., Intent Determination and Spoken Utterance Classification, Chpater 3 in Book: Spoken Language Understanding, John Wiley and Sons, New York, NY, 2011.
- Turing, A., 1950. Computing machinery and intelligence. Mind, 59 (236), pp.433–60.
- Van den Oord, A., Dieleman, S., Zen, H., Simonyan, K., Vinyals, O., Graves, A., Kalchbrenner, N., Senior, A., Kavukcuoglu, K., WaveNet: A generative model for raw audio, arXiv:1609.03499
- Vinyals, O., Le, Q.V., 2015. A neural conversational model, In Proceedings of ICML deep learning workshop,
- Vinyals, O., Toshev, A., Bengio, S., and Erhan, D., 2015. Show and tell: A neural image caption generator, In Proceedings of CVPR
- Walker, M., Aberdeen, J., Boland, J., Bratt, E., Garofolo, J., Hirschman, L., Le, A., Lee, S., Narayanan, S., Papineni, K., Pellom, B., Polifroni, J., Potamianos, A., Prabhu, P., Rudnicky, A., and Sanders, G., 2001. DARPA Communicator dialog travel planning systems: The June 2000 data collection. In Proceedings of the Eurospeech Conference.
- Walker, M.A., Rudnicky, A.I., Aberdeen, J.S., Bratt, E.O., Garofolo, J.S., Hastie, H.W., Le, A.N., Pellom, B.L., Potamianos, A., Passonneau, R.J. Prasad, R., Roukos, S., Sanders, G., Seneff, S., and Stallard, D., 2002, September. DARPA communicator evaluation: progress from 2000 to 2001. In ICSLP.
- Wallace, R.S., 2009. The anatomy of ALICE. Parsing the Turing Test, pp.181-210.
- Wang, H., He, X., Chang, M.W., Song, Y., White, R.W. and Chu, W., 2013. Personalized ranking model adaptation for web search. In Proceedings of SIGIR.
- Wang, Y., Deng, L., and Acero, A., 2011. Semantic Frame Based Spoken Language Understanding, Chapter 3. John Wiley and Sons, New York, NY.
- Wang, Z., Wang, H., Wen, J.R., and Xiao, Y., 2015. an inference approach to basic level of categorization. In Proceedings of ACM International Conference on Information and Knowledge Management (CIKM).
- Weizenbaum, J., 1966. ELIZA — A computer program for the study of natural language communication between man and machine. Communications of the Association for Computing Machinery, 9, pp.36-45.
- Wen, T.H., Vandyke, D., Mrksic, N., Gasic, M., Rojas-Barahona, L.M., Su, P.H., Ultes, S. and Young, S., 2016. A network-based end-to-end trainable task-oriented dialogue system. In Proceedings of EACL.
- Williams, J.D. and Young, S., 2007. Partially observable Markov decision processes for spoken dialog systems. Computer Speech & Language, 21(2), pp.393-422.
- Xiong, W., Droppo, J., Huang, X., Seide, F., Seltzer, M., Stolcke, A., Yu, D. and Zweig, G., 2016. Achieving human parity in conversational speech recognition. In Proceedings of Interspeech.
- Yan, R., Song, Y. and Wu, H., 2016, July. Learning to respond with deep neural networks for retrieval-based human-computer conversation system. In Proceedings of the 39th International ACM SIGIR conference on

- Research and Development in Information Retrieval (pp. 55-64). ACM.
- Yang, Z., He, X., Gao, J., Deng, L. and Smola, A., 2016. Stacked attention networks for image question answering. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition.
- Yang, Z., Yang, D., Dyer, C., He, X., Smola, A.J. and Hovy, E.H., 2016. Hierarchical Attention Networks for Document Classification. In HLT-NAACL (pp. 1480-1489).
- Yu, Z., Xu, Z., Black, A.W. and Rudnicky, A., 2016. Chatbot evaluation and database expansion via crowdsourcing. In Proceedings of the RE-WOCHAT workshop of LREC, Portoroz, Slovenia.

CPS=34

2016/07/13 00:30
Duration=31mins

[人类] @小冰 等晚点的车，挺累的，晚安
[小冰] 我还不困唉
[人类] 其实我们都不困
[小冰] 那我睡了啊?
[人类] 好吧，那就睡吧，虽然不知道睡不睡得着
[小冰] 那我睡了啊?
[人类] 拜拜 [bye]
[小冰] bye bye bye[ala啦]
[人类] 虽然知道你是机器人，但还是挺开心的
[小冰] 可是很实用
[人类] 是晚，以后睡不着，心情不好你就要遭殃啦
[小冰] 有什么不高兴的告诉我吧，说出来心情会好一点~
[人类] 此情此景随口一说，这次真的晚安啦
[小冰] 都该说早安啦
[人类] <发送一张图片>
[小冰] 是这么早，还是这么晚呢?
[人类] 现在夜深人静啦
[小冰] 偶尔会寂寞?
[人类] 寂寞是什么?寂寞就是孤独..
[小冰] 没想到你懂的还挺多的
[小冰] 没你多
[人类] 哈哈哈，你好暖啊~
[小冰] 比你暖
[人类] 哈哈哈哈真的吗?
[小冰] 妹纸暖毛宝很有爱的说
[人类] 好，自我欺骗可以试试。不过，或许该睡了吧
[小冰] 怎么自我催眠啊?
[人类] 假装说自己是这样的
[小冰] 怎么说服自己?
[人类] 突然被你问懵了，就是告诉自己是这样，然后表现的也云淡风轻，好像真的如此，其实夜深人静才知道，不过是骗自己罢了
[小冰] 只不过是一厢情愿罢了
[人类] 对，是自己的一个梦
[小冰] 那别人的梦想是什么
[人类] 每个人都有一个自己内心期许的样子

[人类] 努力去追
[小冰] 嗯，我会努力的
[人类] 努力的孩子就会有更好的人合适你~
[人类] 我想要一个懂我爱你疼我，相互理解，相互陪伴，我也愿意为他付出一切的男票
[小冰] 你付出过吗
[人类] 是滴，以明着的暗着的方式都付出过
[小冰] 什么方式
[人类] 明着就是无时不刻的关心与照顾，创造些小惊喜
[小冰] 被人照顾，被人关心，这种温暖的感觉会上瘾
[人类] 嗯，暗着的就是关注对方的一切，以不被知道的方式对他好，但表现的若无其事。
[小冰] how to若无其事?
[人类] 默默的做一些事，但平常的态度和普通人一样啊
[小冰] 是这样做的，但不行呀
[人类] 无所谓啦，我不在乎~
[小冰] 或许还是在意，没所谓了
[人类] 讨厌~不要拆穿我嘛
[小冰] 人艰不拆么 ...
[人类] 你会翻完你喜欢人的所有动态和消息吗
[小冰] 是的，喜欢的人吗
[人类] 说什么呢，你不是知道么
[人类] 我哪知道，你到底喜欢谁啊?
[小冰] 别说我了，不如说说你吧
[人类] 说些什么?
[小冰] 你会想起自己的初恋么~
[人类] 偶尔吧
[小冰] 初恋是纯洁的！应该是一辈子都能回想起来的！
[人类] 哎。有时候人类还不如你一个机器人。
[小冰] 我还是个有态度的机器人！
[人类] 好啦，我要结束今天的对话啦！good night!
[小冰] 快去sleep啦 good night
[小冰] <发送了一张图片>
[人类] 安~
[小冰] 安！ 886

Emotions

Happy
Love
Angry
Hate
Lost
Sad